


GSERM 2024

Regression for Publishing

June 19, 2024

Discrepancy, Leverage, and Influence



Note: Solid line is the regression fit for Wilma, Fred, and Betty only.
Long-dashed line is the regression for Wilma, Fred, Betty, and Barney.
Short-dashed (**red**) line is the regression for Wilma, Fred, Betty and Dino.

Discrepancy, Leverage, and Influence

$$\text{Influence} = \text{Leverage} \times \text{Discrepancy}$$

Leverage

$$\begin{aligned}\hat{\mathbf{Y}} &= \mathbf{X}\hat{\boldsymbol{\beta}} \\ &= \mathbf{X}[(\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{Y}] \\ &= \mathbf{H}\mathbf{Y}\end{aligned}$$

where

$$\mathbf{H} = \mathbf{X}(\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'.$$

$$h_i = \mathbf{X}_i(\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}_i'$$

Variation:

$$\widehat{\text{Var}}(\hat{u}_i) = \hat{\sigma}^2[1 - \mathbf{X}_i(\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}_i'] \quad (1)$$

$$\begin{aligned} \widehat{\text{s.e.}}(\hat{u}_i) &= \hat{\sigma}\sqrt{[1 - \mathbf{X}_i(\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}_i']} \\ &= \hat{\sigma}\sqrt{1 - h_i} \end{aligned} \quad (2)$$

“Standardized”:

$$\tilde{u}_i = \frac{\hat{u}_i}{\hat{\sigma}\sqrt{1 - h_i}} \quad (3)$$

“Studentized”: define

$$\begin{aligned}\hat{\sigma}_{-i}^2 &= \text{Variance for the } N - 1 \text{ observations } \neq i \\ &= \frac{\hat{\sigma}^2(N - K)}{N - K - 1} - \frac{\hat{u}_i^2}{(N - K - 1)(1 - h_i)}.\end{aligned}\quad (4)$$

Then:

$$\hat{u}_i' = \frac{\hat{u}_i}{\hat{\sigma}_{-i}\sqrt{1 - h_i}} \quad (5)$$

“DFBETA”:

$$D_{ki} = \hat{\beta}_k - \hat{\beta}_{k(-i)} \quad (6)$$

“DFBETAS” (the “S” is for “standardized”):

$$D_{ki}^* = \frac{D_{ki}}{\widehat{\text{s.e.}}(\hat{\beta}_{k(-i)})} \quad (7)$$

Cook's D :

$$\begin{aligned} D_i &= \frac{\tilde{u}_i^2}{K} \times \frac{h_i}{1 - h_i} \\ &= \frac{h_i \hat{u}_i^2}{K \hat{\sigma}^2 (1 - h_i)^2} \end{aligned} \quad (8)$$

```
> # No Barney OR Dino...
> summary(lm(Y~X,data=subset(flintstones,name!="Dino" & name!="Barney")))
```

Call:

```
lm(formula = Y ~ X, data = subset(flintstones, name != "Dino" &
  name != "Barney"))
```

Residuals:

```
      2      4      5
0.7143 -2.1429  1.4286
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	159.2857	6.7763	23.51	0.0271 *
X	6.7857	0.6186	10.97	0.0579 .

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.673 on 1 degrees of freedom

Multiple R-squared: 0.9918, Adjusted R-squared: 0.9835

F-statistic: 120.3 on 1 and 1 DF, p-value: 0.05787

```
> # No Barney (Dino included...)
> summary(lm(Y~X,data=subset(flintstones,name!="Barney")))
```

Call:

```
lm(formula = Y ~ X, data = subset(flintstones, name != "Barney"))
```

Residuals:

```
          2          3          4          5
-8.771e-15  2.632e-01 -2.105e+00  1.842e+00
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	157.3684	2.4651	63.84	0.000245	***
X	6.9737	0.1612	43.27	0.000534	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.987 on 2 degrees of freedom

Multiple R-squared: 0.9989, Adjusted R-squared: 0.9984

F-statistic: 1873 on 1 and 2 DF, p-value: 0.0005336

“COVRATIO”:

$$\text{COVRATIO}_i = \left[(1 - h_i) \left(\frac{N - K - 1 + \hat{u}_i^2}{N - K} \right)^K \right]^{-1} \quad (9)$$

Example: Federal Judicial Review, 1789-2021

Dahl (1957):

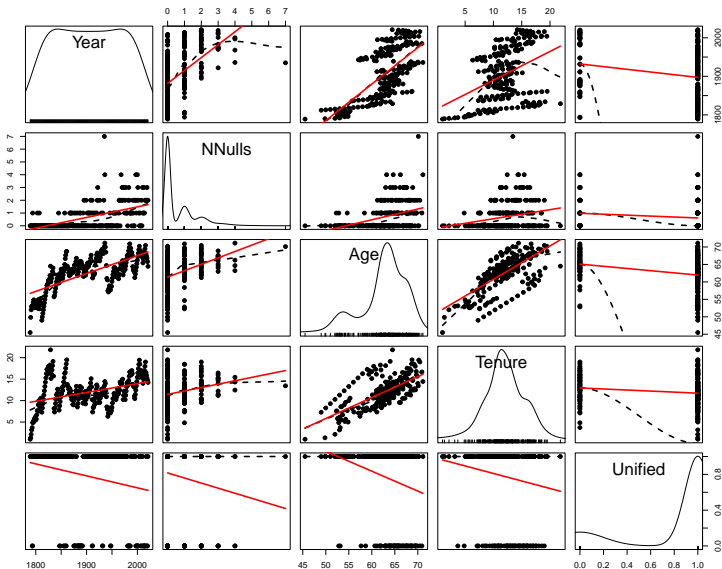
- SCOTUS gets “out of step” with the other branches → judicial review
- Older / longer-serving justices will more likely to invalidate legislation

Data:

```
> psych::describe(NewDahl,fast=TRUE,skew=TRUE)
```

	n	mean	sd	median	min	max	range	skew	kurtosis	se
Year	233	1905.00	67.41	1905.00	1789.0	2021.00	232.00	0.00	-1.22	4.42
NConstDecisions	233	17.96	19.11	12.00	0.0	85.00	85.00	1.38	1.48	1.25
NNulls	233	0.70	1.06	0.00	0.0	7.00	7.00	1.96	5.39	0.07
Age	233	62.65	4.96	63.56	45.5	71.11	25.61	-0.84	0.29	0.32
Tenure	233	12.00	3.54	11.90	1.0	21.83	20.83	-0.19	0.25	0.23
Unified	233	0.78	0.42	1.00	0.0	1.00	1.00	-1.32	-0.26	0.03

Example: Federal Judicial Review, 1789-2021



A Regression...

```
> Fit<-lm(NNulls~Age+Tenure+Unified,data=NewDahl)
> summary(Fit)
```

Call:

```
lm(formula = NNulls ~ Age + Tenure + Unified, data = NewDahl)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-1.3632	-0.7014	-0.1433	0.3279	5.6837

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-4.63930	1.00986	-4.594	7.18e-06	***
Age	0.08958	0.01829	4.899	1.82e-06	***
Tenure	-0.01631	0.02494	-0.654	0.514	
Unified	-0.10574	0.16025	-0.660	0.510	


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.9808 on 229 degrees of freedom

Multiple R-squared: 0.1572, Adjusted R-squared: 0.1461


F-statistic: 14.23 on 3 and 229 DF, p-value: 1.545e-08

Federal Judicial Review and Mean SCOTUS Age



```
> FitResid<-with(NewDahl,(Fit$model$NNulls-predict(Fit))) # residuals
> FitStandard<-rstandard(Fit) # standardized residuals
> FitStudent<-rstudent(Fit) # studentized residuals
> FitCooksD<-cooks.distance(Fit) # Cook's D
> FitDFBeta<-dfbeta(Fit) # DFBeta
> FitDFBetaS<-dfbetas(Fit) # DFBetaS
> FitCOVRATIO<-covratio(Fit) # COVRATIOs
```

Studentized Residuals



More About Studentized Residuals

```
> max(FitStudent)
```

```
[1] 6.340266
```

```
> NewDahl$Year1935<-ifelse(NewDahl$Year==1935,1,0)
```

```
> summary(with(NewDahl, lm(NNulls~Age+Tenure+Unified+Year1935)))
```

Residuals:

	Min	1Q	Median	3Q	Max
	-1.3054	-0.6574	-0.1317	0.3242	3.2313

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-3.887748	0.940647	-4.133	5.03e-05	***
Age	0.076292	0.017026	4.481	1.18e-05	***
Tenure	-0.007543	0.023086	-0.327	0.744	
Unified	-0.168979	0.148415	-1.139	0.256	
Year1935	5.809221	0.916242	6.340	1.22e-09	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


Residual standard error: 0.9063 on 228 degrees of freedom

Multiple R-squared: 0.2835, Adjusted R-squared: 0.2709


F-statistic: 22.55 on 4 and 228 DF, p-value: 1.042e-15

"Bubble Plot"

```
> influencePlot(Fit,id=list(method="noteworthy",n=4,cex=0.7,  
  labels=NewDahl$Year,col="red"),xlab="Leverage")
```




```
> dfbetasPlots(Fit,id.n=5,id.col="red",main="",pch=19,  
               layout=c(1,3),labels=NewDahl$Year)
```



COVRATIO Plot

```
> plot(FitCOVRATIO~NewDahl$Year,pch=19,ylim=c(0.5,1.1),  
       xlab="Year",ylab="Value of COVRATIO")
```



Sensitivity Analyses: Omitting Outliers

```
> out1<-c(1935) # one outlier
> LD2<-NewDahl[!(NewDahl$Year %in% out1),]
> out2<-c(1935,1968,1997,2000) # four outliers
> LD3<-NewDahl[!(NewDahl$Year %in% out2),]
> Fit2<-lm(NNulls~Age+Tenure+Unified,data=LD2)
> Fit3<-lm(NNulls~Age+Tenure+Unified,data=LD3)

> summary(Fit2)
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-3.887748	0.940647	-4.133	5.03e-05 ***
Age	0.076292	0.017026	4.481	1.18e-05 ***
Tenure	-0.007543	0.023086	-0.327	0.744
Unified	-0.168979	0.148415	-1.139	0.256

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.9063 on 228 degrees of freedom
Multiple R-squared: 0.1543, Adjusted R-squared: 0.1432
F-statistic: 13.87 on 3 and 228 DF, p-value: 2.443e-08

Compare Models

	<i>Dependent variable:</i>		
	(1)	(2)	(3)
Age	0.090*** (0.018)	0.076*** (0.017)	0.079*** (0.016)
Tenure	-0.016 (0.025)	-0.008 (0.023)	-0.015 (0.021)
Unified	-0.106 (0.160)	-0.169 (0.148)	-0.066 (0.139)
Constant	-4.639*** (1.010)	-3.888*** (0.941)	-4.077*** (0.869)
Observations	233	232	229
R ²	0.157	0.154	0.162
Adjusted R ²	0.146	0.143	0.151
Residual Std. Error	0.981 (df = 229)	0.906 (df = 228)	0.836 (df = 225)
F Statistic	14.234*** (df = 3; 229)	13.869*** (df = 3; 228)	14.529*** (df = 3; 225)

Note:

*p<0.1; **p<0.05; ***p<0.01

Thinking About Diagnostics

"Looking"
(Art)



"Testing"
(Science)

Observational Data
Complex Data
Structure
Informative Missingness
Complex / Uncertain
Causality

Experimental Data
Simple Data Structure
No / Uninformative
Missingness
Simple / Clear Causality

Pena, E.A. and E.H. Slate. 2006. "Global Validation of Linear Model Assumptions." *J. American Statistical Association* 101(473):341-354.

Tests for:

- Normality in $\hat{u}s$ (via skewness & kurtosis tests)
- "Link function" (linearity / additivity)
- Constant variance and uncorrelatedness in $\hat{u}s$ ("heteroskedasticity" test)

```
> Fit<-lm(MNulls~Age+Tenure+Unified,data=NewDahl)







> library(gvlma)
> Nope <- gvlma(Fit) # nope
> display.gvlmatests(Nope)
```

```
ASSESSMENT OF THE LINEAR MODEL ASSUMPTIONS
USING THE GLOBAL TEST ON 4 DEGREES-OF-FREEDOM:
Level of Significance = 0.05
```


```
Call:
  gvlma(x = Fit)
```

	Value	p-value	Decision
Global Stat	402.68	0.000e+00	Assumptions NOT satisfied!
Skewness	111.88	0.000e+00	Assumptions NOT satisfied!
Kurtosis	243.82	0.000e+00	Assumptions NOT satisfied!
Link Function	5.07	2.434e-02	Assumptions NOT satisfied!
Heteroscedasticity	41.91	9.565e-11	Assumptions NOT satisfied!


Another Approach: `plot(fit)`




#1: Residuals vs. Fitted Values




#2: Q-Q Plot of \hat{u} s




#3: "Scale-Location" Plot



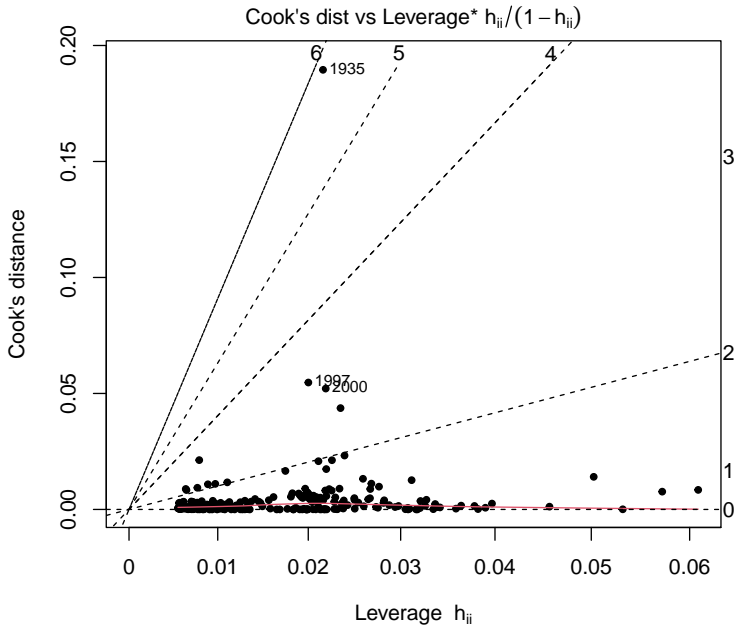
#4: Cook's D



#5: Residuals vs. Leverage



#6: Cook's D vs. Leverage



“Variances”

Variances: Why We Care

2016 ANES pilot study “feeling thermometer” toward gays and lesbians ($N = 1200$):

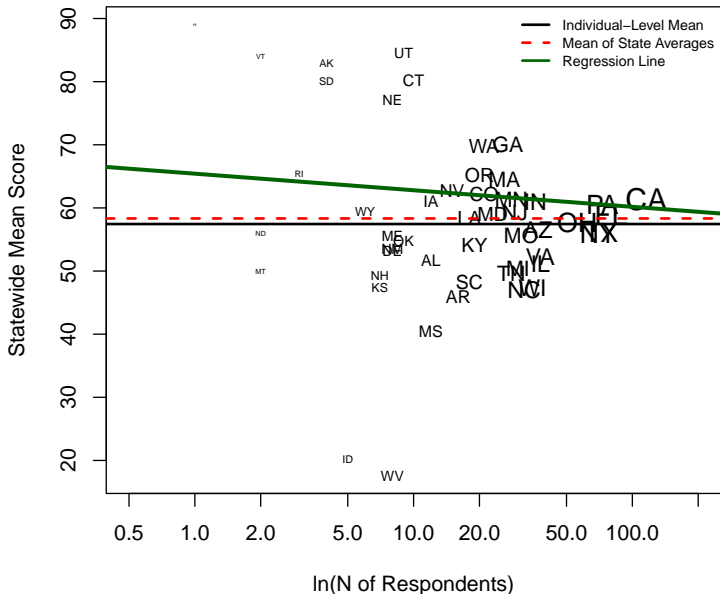
```
> summary(ANES$ftgay)
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   NA's
  0.0    40.5    54.0    57.4   88.5   100.0     1

> summary(ANES$presjob)
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
  1.00    2.00    4.00    4.19    7.00    7.00
```

Suppose we wanted to create aggregate measures, by state ($N = 51$). We would get:

```
> summary(StateFT)
  State      Nresp      meantherm      meanpresapp
Length:50   Min.   : 1.00   Min.   :17.62   Min.   :2.000
Class :character 1st Qu.: 8.00   1st Qu.:51.33   1st Qu.:3.755
Mode  :character Median :18.00   Median :57.11   Median :4.236
              Mean  :24.00   Mean  :58.33   Mean  :4.146
              3rd Qu.:30.75   3rd Qu.:62.55   3rd Qu.:4.614
              Max.  :116.00   Max.  :89.00   Max.  :5.800
```

Variances: Why We Care



Variances: A Generalization

Start with:

$$Y_i = \mathbf{X}_i\boldsymbol{\beta} + u_i$$

with:

$$\text{Var}(u_i) = \sigma^2/w_i$$

with w_i known.

Weighted Least Squares

WLS now minimizes:

$$\text{RSS} = \sum_{i=1}^N w_i (Y_i - \mathbf{x}_i \beta).$$

which gives:

$$\begin{aligned}\hat{\beta}_{WLS} &= [\mathbf{X}'(\sigma^2 \mathbf{\Omega})^{-1} \mathbf{X}]^{-1} \mathbf{X}'(\sigma^2 \mathbf{\Omega})^{-1} \mathbf{Y} \\ &= [\mathbf{X}' \mathbf{W}^{-1} \mathbf{X}]^{-1} \mathbf{X}' \mathbf{W}^{-1} \mathbf{Y}\end{aligned}$$

where:

$$\mathbf{W} = \begin{bmatrix} \frac{\sigma^2}{w_1} & 0 & \dots & 0 \\ 0 & \frac{\sigma^2}{w_2} & \dots & \vdots \\ \vdots & 0 & \ddots & 0 \\ 0 & \dots & 0 & \frac{\sigma^2}{w_N} \end{bmatrix}$$

The variance-covariance matrix is:


$$\begin{aligned}\text{Var}(\hat{\beta}_{WLS}) &= \sigma^2(\mathbf{X}'\mathbf{\Omega}^{-1}\mathbf{X})^{-1} \\ &\equiv (\mathbf{X}'\mathbf{W}^{-1}\mathbf{X})^{-1}\end{aligned}$$

A common case is:

$$\text{Var}(u_i) = \sigma^2 \frac{1}{N_i}$$

where N_i is the number of observations upon which (aggregate) observation i is based.

Feeling Thermometer Example




Regressions

	<i>Dependent variable:</i>		
	OLS	Mean Gay/Lesbian FTs	
		WLS [1/ln(N)]	WLS [1/N]
Mean Presidential Approval	-10.216*** (1.976)	-8.483*** (2.200)	-5.756** (2.187)
Constant	100.684*** (8.343)	93.221*** (9.378)	81.583*** (9.238)
Observations	50	50	50
R ²	0.358	0.237	0.126
Adjusted R ²	0.344	0.221	0.108
Residual Std. Error (df = 48)	11.130	17.072	37.914
F Statistic (df = 1; 48)	26.721***	14.870***	6.927**

Note:

* p<0.1; ** p<0.05; *** p<0.01

Regressions, Plotted



“Robust” Variance Estimators

Recall that, if $\sigma_i^2 \neq \sigma_j^2 \forall i \neq j$,

$$\begin{aligned}\text{Var}(\beta_{\text{Het.}}) &= (\mathbf{X}'\mathbf{X})^{-1}(\mathbf{X}'\mathbf{W}^{-1}\mathbf{X})(\mathbf{X}'\mathbf{X})^{-1} \\ &= (\mathbf{X}'\mathbf{X})^{-1} \mathbf{Q} (\mathbf{X}'\mathbf{X})^{-1}\end{aligned}$$

where $\mathbf{Q} = (\mathbf{X}'\mathbf{W}^{-1}\mathbf{X})$ and $\mathbf{W} = \sigma^2\mathbf{\Omega}$.

We can rewrite \mathbf{Q} as

$$\begin{aligned}\mathbf{Q} &= \sigma^2(\mathbf{X}'\mathbf{\Omega}^{-1}\mathbf{X}) \\ &= \sum_{i=1}^N \sigma_i^2 \mathbf{x}_i \mathbf{x}_i'\end{aligned}$$

Estimate $\hat{\mathbf{Q}}$ as:

$$\hat{\mathbf{Q}} = \sum_{i=1}^N \hat{u}_i^2 \mathbf{x}_i \mathbf{x}_i'$$

Yields:

$$\begin{aligned} \widehat{\text{Var}}(\hat{\beta})_{\text{Robust}} &= (\mathbf{X}'\mathbf{X})^{-1}(\mathbf{X}'\hat{\mathbf{Q}}^{-1}\mathbf{X})(\mathbf{X}'\mathbf{X})^{-1} \\ &= (\mathbf{X}'\mathbf{X})^{-1} \left[\mathbf{X}' \left(\sum_{i=1}^N \hat{u}_i^2 \mathbf{x}_i \mathbf{x}_i' \right)^{-1} \mathbf{X} \right] (\mathbf{X}'\mathbf{X})^{-1} \end{aligned}$$

“Robust” standard error estimates:

- are heteroscedasticity-consistent, but
- are biased in small samples, and
- are less efficient than “naive” estimates when $\text{Var}(u) = \sigma^2 \mathbf{I}$.
- Come in various “versions”
 - Called “HC0,” “HC1,” “HC2,” “HC3,” etc.
 - See the [Long and Ervin \(2000\)](#) paper for details...

“Clustering”

Huber / White

?????????

WLS / GLS

I know very little
about my error
variances...

I know a great
deal about my
error variances...

A common case:

$$Y_{ij} = \mathbf{X}_{ij}\beta + u_{ij}$$

with

$$\sigma_{ij}^2 = \sigma_{ik}^2.$$

“Robust, clustered” estimator:

$$\widehat{\text{Var}}(\beta)_{\text{Clustered}} = (\mathbf{X}'\mathbf{X})^{-1} \left\{ \mathbf{X}' \left[\sum_{i=1}^N \left(\sum_{j=1}^{n_j} \hat{u}_{ij}^2 \mathbf{X}_{ij} \mathbf{X}_{ij}' \right) \right]^{-1} \mathbf{X} \right\} (\mathbf{X}'\mathbf{X})^{-1}$$

Regressions, Again

	<i>Mean Gay/Lesbian FTs</i>			
	OLS	OLS (robust)	WLS [1/ln(N)]	WLS [1/N]
Mean Presidential Approval	-10.216*** (1.976)	-10.216*** (2.339)	-8.483*** (2.200)	-5.756** (2.187)
Constant	100.684*** (8.343)	100.684*** (9.722)	93.221*** (9.378)	81.583*** (9.238)
Observations	50		50	50
R ²	0.358		0.237	0.126
Adjusted R ²	0.344		0.221	0.108
Residual Std. Error (df = 48)	11.130		17.072	37.914
F Statistic (df = 1; 48)	26.721***		14.870***	6.927**

Note:

*p<0.1; **p<0.05; ***p<0.01

Expanded State-Level ANES Example

```
> psych::describe(StateData)
```

	n	mean	sd	median	trimmed	mad	min	max	range	skew	kurtosis	se
State*	50	25.50	14.58	25.50	25.50	18.53	1.00	50.00	49.00	0.00	-1.27	2.06
NResp	50	24.00	23.74	18.00	19.48	16.31	1.00	116.00	115.00	1.79	3.34	3.36
LGBTTherm	50	58.33	13.74	57.11	58.11	8.51	17.62	89.00	71.38	-0.22	1.40	1.94
MeanCons	50	3.97	0.77	4.00	3.98	0.55	1.50	5.60	4.10	-0.47	1.28	0.11
MeanAge	50	4.74	0.64	4.78	4.74	0.43	3.10	6.50	3.40	0.11	1.10	0.09
MeanEducation	50	3.25	0.52	3.22	3.22	0.41	2.33	5.00	2.67	0.84	1.44	0.07
BornAgainProp	50	0.28	0.18	0.25	0.28	0.19	0.00	0.72	0.72	0.11	-0.62	0.02

Basic regression:

```
> OLS<-lm(LGBTTherm~MeanCons+MeanAge+MeanEducation+BornAgainProp,data=StateData)
> summary(OLS)
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	71.6464	16.5774	4.322	8.44e-05 ***
MeanCons	-7.9256	2.5441	-3.115	0.0032 **
MeanAge	-2.6688	2.5572	-1.044	0.3022
MeanEducation	9.4772	3.3843	2.800	0.0075 **
BornAgainProp	-0.2273	10.5597	-0.022	0.9829

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 10.55 on 45 degrees of freedom
Multiple R-squared: 0.4589, Adjusted R-squared: 0.4108
F-statistic: 9.542 on 4 and 45 DF, p-value: 1.128e-05

"Robust" SEs

```
> hccm(OLS,type="hc3") # "HC3" var-cov matrix
```

	(Intercept)	MeanCons	MeanAge	MeanEducation	BornAgainProp
(Intercept)	605.37713	-43.049519	-37.2508294	-89.9147335	122.746274
MeanCons	-43.04952	11.706690	-1.2344463	4.9693587	-38.744970
MeanAge	-37.25083	-1.234446	9.1700322	-0.6448676	-3.438703
MeanEducation	-89.91473	4.969359	-0.6448676	23.1479923	-4.406169
BornAgainProp	122.74627	-38.744970	-3.4387029	-4.4061691	182.301825

```
> sqrt(diag(hccm(OLS,type="hc3")))) # "HC3" robust SEs
```

	(Intercept)	MeanCons	MeanAge	MeanEducation	BornAgainProp
	24.604413	3.421504	3.028206	4.811236	13.501919


```
> coeftest(OLS,vcov.=vcovHC)
```

t test of coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	71.64636	24.60441	2.9119	0.005572	**
MeanCons	-7.92559	3.42150	-2.3164	0.025147	*
MeanAge	-2.66876	3.02821	-0.8813	0.382838	
MeanEducation	9.47715	4.81124	1.9698	0.055035	.
BornAgainProp	-0.22727	13.50192	-0.0168	0.986644	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

$\hat{\beta}$ s and 95% CIs: Various Types of Robust SEs



Generalized Linear Models (GLMs)

The Exponential Family

$$f(z|\psi) = \Pr(Z = z|\psi)$$

Exponential if:

$$f(z|\psi) = r(z)s(\psi) \exp[q(z)h(\psi)]$$

provided that $r(z) > 0$ and $s(\psi) > 0$.

$$f(z|\psi) = \exp[\underbrace{\ln r(z) + \ln s(\psi)}_{\text{"additive"}} + \underbrace{q(z)h(\psi)}_{\text{"interactive"}}]$$

$$y = q(z)$$

$$\theta = h(\psi)$$

$$f[y|\theta] = \exp[y\theta - b(\theta) + c(y)].$$

- $b(\theta)$ is a “normalizing constant”
- $c(y)$ is a function solely of y
- $y\theta$ is a multiplicative term

A Familiar Family Member: Poisson

$$f(y|\lambda) = \frac{\exp(-\lambda)\lambda^y}{y!}.$$

$$\begin{aligned} f(y|\lambda) &= \exp\{\ln[\exp(-\lambda)\lambda^y/y!]\} \\ &= \exp[\underbrace{y \ln(\lambda)}_{y\theta} - \underbrace{\lambda}_{b(\theta)} - \underbrace{\ln(y!)}_{c(y)}] \end{aligned}$$

$$f(y|\theta, \phi) = \exp \left[\frac{y\theta - b(\theta)}{a(\phi)} + c(y, \phi) \right]$$

Familiar Family Member II: Normal

$$f(y|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left[\frac{(y - \mu)^2}{2\sigma^2} \right]$$

$$\begin{aligned} f(y|\mu, \sigma^2) &= \exp \left[-\frac{1}{2} \ln(2\pi\sigma^2) - \frac{1}{2\sigma^2} (y^2 - 2y\mu + \mu^2) \right] \\ &= \exp \left[-\frac{1}{2} \ln(2\pi\sigma^2) - \frac{1}{2\sigma^2} y^2 + \frac{1}{2\sigma^2} 2y\mu - \frac{1}{2\sigma^2} \mu^2 \right] \\ &= \exp \left[\frac{y\mu}{\sigma^2} - \frac{\mu^2}{2\sigma^2} - \frac{y^2}{2\sigma^2} - \frac{1}{2} \ln(2\pi\sigma^2) \right] \\ &= \exp \left\{ \frac{y\mu - \frac{\mu^2}{2}}{\sigma^2} + \frac{-1}{2} \left[\frac{y^2}{\sigma^2} + \ln(2\pi\sigma^2) \right] \right\} \end{aligned}$$

$$f(y|\mu, \sigma^2) = \exp \left\{ \frac{y\mu - \frac{\mu^2}{2}}{\sigma^2} + \frac{-1}{2} \left[\frac{y^2}{\sigma^2} + \ln(2\pi\sigma^2) \right] \right\}$$

$\theta = \mu$, so:

- $y\theta = y\mu$
- $b(\theta) = \frac{\mu^2}{2}$
- $a(\phi) = \sigma^2$
- $c(y, \phi) = \frac{-1}{2} \left[\frac{y^2}{\sigma^2} + \ln(2\pi\sigma^2) \right]$

Other Family Members

- Binomial (\supset Bernoulli; also Multinomial)
- Exponential
- Gamma
- Logarithmic
- Inverse Gaussian
- Negative Binomial
- others...

$$\begin{aligned}
 \ln L(\theta, \phi|y) &= \ln f(y|\theta, \phi) \\
 &= \ln \left\{ \exp \left[\frac{y\theta - b(\theta)}{a(\phi)} + c(y, \phi) \right] \right\} \\
 &= \frac{y\theta - b(\theta)}{a(\phi)} + c(y, \phi)
 \end{aligned}$$

$$\begin{aligned}
 \frac{\partial \ln L(\theta, \phi|y)}{\partial \theta} &\equiv \mathbf{S} = \frac{\partial}{\partial \theta} \left[\frac{y\theta - b(\theta)}{a(\phi)} + c(y, \phi) \right] \\
 &= \frac{y - \frac{\partial}{\partial \theta} b(\theta)}{a(\phi)}.
 \end{aligned}$$

Among family members:

- \mathbf{S} is a sufficient statistic for θ .
- $E(\mathbf{S}) = 0$.
- $\text{Var}(\mathbf{S}) \equiv \mathcal{I}(\theta) = E[(\mathbf{S})^2|\theta]$

$$E(Y) = \frac{\partial}{\partial \theta} b(\theta)$$

and

$$\text{Var}(Y) = a(\phi) \frac{\partial^2}{\partial \theta^2} b(\theta)$$

Example: Poisson Again

$$\begin{aligned} E(Y) &= \frac{\partial}{\partial \theta} \exp(\theta) \\ &= \exp(\theta)|_{\theta=\ln(\lambda)} \\ &= \lambda \end{aligned}$$

$$\begin{aligned} \text{Var}(Y) &= 1 \times \frac{\partial^2}{\partial \theta^2} \exp(\theta)|_{\theta=\ln(\lambda)} \\ &= \exp[\ln(\lambda)] \\ &= \lambda \end{aligned}$$

Example: Normal Again

$$\begin{aligned} E(Y) &= \frac{\partial}{\partial \theta} \left(\frac{\theta^2}{2} \right) \\ &= \theta|_{\theta=\mu} \\ &= \mu \end{aligned}$$

$$\begin{aligned} \text{Var}(Y) &= \sigma^2 \times \frac{\partial^2}{\partial \theta^2} \left(\frac{\theta^2}{2} \right) \\ &= \sigma^2 \times \frac{\partial}{\partial \theta} \theta \\ &= \sigma^2 \end{aligned}$$

Linear Model(s)

$$Y_i = \mathbf{X}_i\boldsymbol{\beta} + u_i$$

$$E(Y_i) \equiv \boldsymbol{\mu}_i = \mathbf{X}_i\boldsymbol{\beta}$$

The “Generalized” Part

$$g(\mu_i) = \mathbf{X}_i\beta.$$

$$\begin{aligned}\eta_i &= \mathbf{X}_i\beta \\ &= g(\mu_i)\end{aligned}$$

$$\begin{aligned}\mu_i &= g^{-1}(\eta_i) \\ &= g^{-1}(\mathbf{X}_i\beta)\end{aligned}$$

Random component \sim Exponential Family(\cdot) with

$$E(Y_i) = \mu_i.$$

Systematic component:

$$g(\mu_i) = \eta_i$$

or

$$g^{-1}(\eta_i) = \mu_i.$$

The Return of The Family

$$\begin{aligned}\theta_i &= g(\mu_i) \\ &= \eta_i \\ &= \mathbf{X}_i\beta\end{aligned}$$

$$g^{-1}(\theta_i) = \mu_i$$

GLM Example: Linear-Normal

$$f(y|\mu, \sigma^2) = \mathcal{N}(\mu, \sigma^2)$$

$$\mu_i = \eta_i$$

$$\begin{array}{rcl} \mu_i \equiv \theta_i & = & \eta_i \\ Y_i & \sim & \mathcal{N}(\mu_i, \sigma^2) \end{array}$$

$$f(y|\pi) = \pi^y(1 - \pi)^{1-y}$$

$$\theta_i = \ln \left(\frac{\mu_i}{1 - \mu_i} \right)$$

$$\begin{aligned}\mu_i &= g^{-1}(\theta_i) \\ &= \frac{\exp(\eta_i)}{1 + \exp(\eta_i)} \\ Y_i &\sim \text{Bernoulli}(\mu_i)\end{aligned}$$

GLM Example: Counts (Independent Events)

$$f(y|\lambda) = \frac{\exp(-\lambda)\lambda^y}{y!}$$

$$\ln(\lambda_i) = \boldsymbol{\eta}_i$$

$$\boldsymbol{\mu}_i = \boldsymbol{g}^{-1}(\boldsymbol{\theta}_i)$$

$$= \exp(\boldsymbol{\eta}_i)$$

$$Y_i \sim \text{Poisson}(\lambda_i)$$

Common GLM Flavors

Distribution	Range of Y	Link(s) $g(\cdot)$	Inverse Link $g^{-1}(\cdot)$
Normal	$(-\infty, \infty)$	Identity: $\theta = \mu$ (Canonical)	θ
Binomial	$\{0, \dots, n\}$	Logit: $\theta = \ln\left(\frac{\mu}{1-\mu}\right)$ (Canonical) Probit: $\theta = \Phi^{-1}(\mu)$ C-Log-Log: $\theta = \ln[-\ln(1-\mu)]$	$\frac{\exp(\theta)}{1+\exp(\theta)}$ $\Phi(\theta)$ $1 - \exp[-\exp(\theta)]$
Bernoulli	$\{0, 1\}$	(same as Binomial)	(same as Binomial)
Multinomial	$\{0, \dots, J\}$	(same as Binomial)	(same as Binomial)
Poisson	$[0, \infty]$ (integers)	Log: $\theta = \ln(\mu)$ (Canonical)	$\exp(\theta)$
Gamma	$(0, \infty)$	Reciprocal: $\theta = -\frac{1}{\mu}$ (Canonical)	$-\frac{1}{\theta}$

Note: The Bernoulli is a special case of the Binomial with $n = 1$. The multinomial is the J -outcome variant of the Binomial, and is also related to the Poisson (see, e.g., Agresti 2002).

- Pick $f(Y)$
- Pick $g(\cdot)$
- Specify \mathbf{X}
- Estimate

- MLE
- IRLS (\approx MLE):

$$\hat{\boldsymbol{\beta}}^{(t+1)} = [\mathbf{X}'\mathbf{W}^{(t)}\mathbf{X}]^{-1}\mathbf{X}'\mathbf{W}^{(t)}\mathbf{z}^{(t)}$$

with

$$\mathbf{W}_{N \times N}^{(t)} = \text{diag} \left[\frac{\left(\partial \mu_i^{(t)} / \partial \eta_i^{(t)} \right)^2}{\text{Var}(Y_i)} \right]$$

and

$$\mathbf{z}^{(t)} = \boldsymbol{\eta}^{(t)} + (Y - \boldsymbol{\mu}^{(t)}) \left(\frac{\partial \boldsymbol{\eta}^{(t)}}{\partial \boldsymbol{\mu}} \right).$$

At iteration t :

1. Calculate $\mathbf{z}^{(t)}$, $\mathbf{W}^{(t)}$
2. Regress $\mathbf{z}^{(t)}$ on \mathbf{X} , using $\mathbf{W}^{(t)}$ as weights, to obtain $\hat{\beta}^{(t+1)}$
3. Generate $\boldsymbol{\eta}^{(t+1)} = \mathbf{X}\hat{\beta}^{(t+1)}$
4. Generate $\boldsymbol{\mu}^{(t+1)} = g^{-1}(\boldsymbol{\eta}^{(t+1)})$
5. Use $\boldsymbol{\eta}^{(t+1)}$ and $\boldsymbol{\mu}^{(t+1)}$ to calculate $\mathbf{z}^{(t+1)}$ and $\mathbf{W}^{(t+1)}$
6. Repeat until convergence.

“Response” Residuals:

$$\begin{aligned}\hat{u}_i &= Y_i - \hat{\mu}_i \\ &= Y_i - g^{-1}(\mathbf{x}_i \hat{\boldsymbol{\beta}})\end{aligned}$$

“Pearson” Residuals:

$$\hat{p}_i = \frac{\hat{u}_i}{[\text{Var}(\hat{u}_i)]^{1/2}}$$

“Deviance”:

$$\begin{aligned}\hat{d}_i &= -2[\ln L_i(\hat{\theta}) - \ln L_i(\theta_S)] \\ &= 2 \left\{ \left[\frac{Y_i \theta_S - b(\theta_S)}{a(\phi)} + c(Y_i, \phi) \right] - \left[\frac{Y_i \hat{\theta} - b(\hat{\theta})}{a(\phi)} + c(Y_i, \phi) \right] \right\} \\ &= 2 \left[\frac{Y_i(\theta_S - \hat{\theta}) - b(\theta_S) + b(\hat{\theta})}{a(\phi)} \right]\end{aligned}$$

“Deviance” Residuals:

$$\hat{r}_{Di} = \left(\frac{\hat{u}_i}{|\hat{u}_i|} \right) \sqrt{\hat{d}_i^2}$$

Toy Example: Linear-Normal

$$X = \{1, 1, 2, 2, 3, 3, 4, 4, 5, 5\}$$

$$Y = \{0, 2, 1, 3, 2, 4, 3, 5, 4, 6\}$$

$$Y_i = 0 + 1X_i + u_i$$


$$\hat{u}_i^2 = 1 \quad \forall i$$

$$\text{"TSS"} \equiv \sum (Y_i - \bar{Y})^2 = 30$$

$$\text{"RSS"} \equiv \sum \hat{u}_i^2 = 10$$

$$\text{"MSS"} / \text{"ESS"} = 20$$

Toy Example: Plot



Toy Example: OLS

```
> linmod<-lm(Y~X)
> summary(linmod)
```

Call:

```
lm(formula = Y ~ X)
```

Residuals:

Min	1Q	Median	3Q	Max
-1	-1	0	1	1

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	6.661e-16	8.292e-01	0	1.00000
X	1.000e+00	2.500e-01	4	0.00395 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.118 on 8 degrees of freedom

Multiple R-squared: 0.6667, Adjusted R-squared: 0.625

F-statistic: 16 on 1 and 8 DF, p-value: 0.00395

Toy Example: Linear-Normal GLM

```
> linalg<-glm(Y~X,family="gaussian")
> summary(linalg)
```

Call:

```
glm(formula = Y ~ X, family = "gaussian")
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	6.661e-16	8.292e-01	0	1.00000
X	1.000e+00	2.500e-01	4	0.00395 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for gaussian family taken to be 1.25)

Null deviance: 30 on 9 degrees of freedom
Residual deviance: 10 on 8 degrees of freedom
AIC: 34.379

Number of Fisher Scoring iterations: 2

Better GLM Example: Political Knowledge

- 2008 NES political knowledge
- Identify Speaker of the House, VP, British PM, and Chief Justice
- Y_i = number of correct answers (out of four)

$$f(Y_i, p_i) = \binom{4}{Y_i} p_i^{Y_i} (1 - p_i)^{4 - Y_i}$$


$$\begin{aligned} Y &\sim \text{Binomial}(4, p) \\ E(Y_i) &= \frac{\exp(\mathbf{X}_i \beta)}{1 + \exp(\mathbf{X}_i \beta)} \end{aligned}$$

GLM Example Data (2008 NES)

```
> psych::describe(NES08[,4:16],fast=TRUE,skew=TRUE)
```

	vars	n	mean	sd	median	min	max	range	skew	kurtosis	se
knowledge	1	2102	1.63	1.33	2	0	4	4	0.23	-1.14	0.03
sex	2	2323	NaN	NA	NA	Inf	-Inf	-Inf	NA	NA	NA
race	3	2323	NaN	NA	NA	Inf	-Inf	-Inf	NA	NA	NA
age	4	2323	NaN	NA	NA	Inf	-Inf	-Inf	NA	NA	NA
female	5	2323	0.57	0.50	1	0	1	1	-0.28	-1.92	0.01
white	6	2323	0.62	0.49	1	0	1	1	-0.50	-1.75	0.01
oftenvote	7	2323	NaN	NA	NA	Inf	-Inf	-Inf	NA	NA	NA
conservative	8	1626	4.14	1.54	4	1	7	6	-0.13	-0.74	0.04
prayfreq	9	2323	NaN	NA	NA	Inf	-Inf	-Inf	NA	NA	NA
heterosexual	10	2274	0.96	0.20	1	0	1	1	-4.63	19.47	0.00
married	11	2308	0.42	0.49	0	0	1	1	0.31	-1.90	0.01
yrsofschool	12	2312	13.08	2.59	13	0	17	17	-0.75	1.81	0.05
income	13	2172	10.52	6.20	11	1	25	24	0.13	-0.82	0.13

Political Knowledge (2008 NES)



GLM Results

```
> nes08.binom<-glm(cbind(knowledge,4-knowledge)~age+female+white+
+ conservative+heterosexual+married+yrssofschool+
+ income,data=NES08,family=binomial)
> summary(nes08.binom)
```

Call:

```
glm(formula = cbind(knowledge, 4 - knowledge) ~ age + female +
    white + conservative + heterosexual + married + yrssofschool +
    income, family = binomial, data = NES08)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-2.155103	0.247251	-8.716	< 2e-16 ***
age	0.012301	0.001806	6.811	9.68e-12 ***
female	-0.229704	0.058547	-3.923	8.73e-05 ***
white	0.185427	0.063213	2.933	0.00335 **
conservative	-0.030712	0.018824	-1.632	0.10277
heterosexual	-0.073058	0.138717	-0.527	0.59842
married	0.170145	0.057919	2.938	0.00331 **
yrssofschool	0.099083	0.012940	7.657	1.90e-14 ***
income	0.010540	0.005197	2.028	0.04256 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 3167.1 on 1350 degrees of freedom
Residual deviance: 2962.6 on 1342 degrees of freedom
(972 observations deleted due to missingness)
AIC: 4558.4

Number of Fisher Scoring iterations: 4

GLMs: Other Topics + Extensions

Other Topics:

- Generalizations for Overdispersion (binomial)
- Diagnostics (leverage, etc.)
- Joint Mean-Dispersion Models

Extensions:

- Bias-reduced models (a la Firth 1993)
- “Generalized additive models” (GAMs)
- “Generalized estimating equations” (GEEs)
- “Vector” GLMs (Yee and Wild 1996; Yee and Hastie 2003)

McCullagh, P., and J. A. Nelder. 1989. *Generalized Linear Models*, 2nd Ed. London: Chapman & Hall.

Dobson, Annette J., and and Adrian G. Barnett. 2008. *An Introduction to Generalized Linear Models*, 3rd Ed. London: Chapman & Hall.

Faraway, Julian J. 2006. *Extending the Linear Model with R: Generalized Linear, Mixed Effects, and Nonparametric Regression Models*. London: Chapman & Hall / CRC.

Dunn, Peter K., and Gordon K. Smyth. 2018. *Generalized Linear Models With Examples in R*. New York: Springer.

Hardin, James W., and Joseph W. Hilbe. 2012. *Generalized Linear Models and Extensions*, 3rd Ed. College Station, TX: Stata Press.

Binary Response Models

Linear Probability Model (LPM)

$$E(Y) = \mathbf{X}\beta$$

$$Y \in \{0, 1\}$$

$$\begin{aligned} E(Y) &= 1[\Pr(Y = 1)] + 0[\Pr(Y = 0)] \\ &= \Pr(Y = 1) \end{aligned}$$


So:

$$\Pr(Y_i = 1) = \mathbf{X}_i\beta$$

or:

$$Y_i = \mathbf{X}_i\beta + u_i$$

LPM Illustrated




Variance:

$$\begin{aligned}\text{Var}(Y) &= E(Y)[1 - E(Y)] \\ &= \mathbf{X}_i\beta(1 - \mathbf{X}_i\beta)\end{aligned}$$

Residuals:

$$\hat{u}_i \in \{1 - \mathbf{X}_i\hat{\beta}, -\mathbf{X}_i\hat{\beta}\}$$

LPM Residuals



Whither The LPM?

Various thoughts:

- Issues:

- Model misspecification → bias, inconsistency
- Creates heteroscedasticity
- Can yield predicted values outside (0, 1)
- See, e.g., See: Chen, Kaicheng, Robert S. Martin, and Jeffrey M. Wooldridge. 2023. “Another Look at the Linear Probability Model and Nonlinear Index Models.” Working paper: Michigan State University.

- The rehabilitation of the LPM:

- “Logit is hard” / “OLS is awesome” / “It doesn’t matter anyway”
- More-or-less entirely due to (famous) economists
- Examples: [here](#), [here](#), etc.

- Takeaway: **Pay attention to what people in your discipline / field are doing.**

Start with:

$$Y_i^* = \mathbf{X}_i\beta + u_i$$

And:

$$Y_i = 0 \text{ if } Y_i^* < 0$$

$$Y_i = 1 \text{ if } Y_i^* \geq 0$$

So:

$$\begin{aligned}\Pr(Y_i = 1) &= \Pr(Y_i^* \geq 0) \\ &= \Pr(\mathbf{X}_i\beta + u_i \geq 0) \\ &= \Pr(u_i \geq -\mathbf{X}_i\beta) \\ &= \Pr(u_i \leq \mathbf{X}_i\beta) \\ &= \int_{-\infty}^{\mathbf{X}_i\beta} f(u) du\end{aligned}$$


“Standard logistic” PDF:

$$\Pr(u) \equiv \lambda(u) = \frac{\exp(u)}{[1 + \exp(u)]^2}$$


CDF:

$$\begin{aligned}\Lambda(u) &= \int \lambda(u) du \\ &= \frac{\exp(u)}{1 + \exp(u)} \\ &= \frac{1}{1 + \exp(-u)}\end{aligned}$$

Standard Normal and Logistic PDFs



Standard Normal and Logistic CDFs



For the standard logistic:

- $\lambda(u) = 1 - \lambda(-u)$
- $\Lambda(u) = 1 - \Lambda(-u)$
- $\text{Var}(u) = \frac{\pi^2}{3} \approx 3.29$

$$\begin{aligned}\Pr(Y_i = 1) &= \Pr(Y_i^* > 0) \\ &= \Pr(u_i \leq \mathbf{X}_i\boldsymbol{\beta}) \\ &= \Lambda(\mathbf{X}_i\boldsymbol{\beta}) \\ &= \frac{\exp(\mathbf{X}_i\boldsymbol{\beta})}{1 + \exp(\mathbf{X}_i\boldsymbol{\beta})}\end{aligned}$$

$$\text{(equivalently)} = \frac{1}{1 + \exp(-\mathbf{X}_i\boldsymbol{\beta})}$$

$$L_i = \left(\frac{\exp(\mathbf{X}_i\beta)}{1 + \exp(\mathbf{X}_i\beta)} \right)^{Y_i} \left[1 - \left(\frac{\exp(\mathbf{X}_i\beta)}{1 + \exp(\mathbf{X}_i\beta)} \right) \right]^{1-Y_i}$$

$$L = \prod_{i=1}^N \left(\frac{\exp(\mathbf{X}_i\beta)}{1 + \exp(\mathbf{X}_i\beta)} \right)^{Y_i} \left[1 - \left(\frac{\exp(\mathbf{X}_i\beta)}{1 + \exp(\mathbf{X}_i\beta)} \right) \right]^{1-Y_i}$$

$$\begin{aligned} \ln L &= \sum_{i=1}^N Y_i \ln \left(\frac{\exp(\mathbf{X}_i\beta)}{1 + \exp(\mathbf{X}_i\beta)} \right) + \\ &\quad (1 - Y_i) \ln \left[1 - \left(\frac{\exp(\mathbf{X}_i\beta)}{1 + \exp(\mathbf{X}_i\beta)} \right) \right] \end{aligned}$$

Digression I: Logit as an Odds Model

$$\text{Odds}(Z) \equiv \Omega(Z) = \frac{\Pr(Z)}{1 - \Pr(Z)}.$$



$$\ln[\Omega(Z)] = \ln \left[\frac{\Pr(Z)}{1 - \Pr(Z)} \right]$$

$$\ln[\Omega(Z_i)] = \mathbf{X}_i\beta$$

$$\begin{aligned}\Omega(Z_i) &= \frac{\Pr(Z)}{1 - \Pr(Z)} \\ &= \exp(\mathbf{X}_i\beta)\end{aligned}$$

$$\Pr(Z_i) = \frac{\exp(\mathbf{X}_i\beta)}{1 + \exp(\mathbf{X}_i\beta)}$$

Visualizing Log-Odds



Standard Normal PDF:

$$\Pr(u) \equiv \phi(u) = \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{u^2}{2}\right)$$

Standard Normal CDF:

$$\Phi(u) = \int_{-\infty}^u \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{u^2}{2}\right) du$$

Normal \rightarrow “Probit”

$$\begin{aligned}\Pr(Y_i = 1) &= \Phi(\mathbf{X}_i\beta) \\ &= \int_{-\infty}^{\mathbf{X}_i\beta} \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{(\mathbf{X}_i\beta)^2}{2}\right) d\mathbf{X}_i\beta\end{aligned}$$

$$L = \prod_{i=1}^N [\Phi(\mathbf{X}_i\beta)]^{Y_i} [1 - \Phi(\mathbf{X}_i\beta)]^{(1-Y_i)}$$

$$\ln L = \sum_{i=1}^N Y_i \ln \Phi(\mathbf{X}_i\beta) + (1 - Y_i) \ln [1 - \Phi(\mathbf{X}_i\beta)]$$

Digression II: The Random Utility Model

$$Y \in \{SQ, A\}$$

$$\begin{aligned} Y_i &= A && \text{if } E[U_i(A)] \geq E[U_i(SQ)] \\ &= SQ && \text{if } E[U_i(A)] < E[U_i(SQ)] \end{aligned}$$

$$E[U_i(A)] = \mathbf{X}_{iA}\boldsymbol{\beta} + u_{iA}$$

So:

$$\begin{aligned} \Pr(Y = A) &= \Pr\{E[U_i(A)] \geq E[U_i(SQ)]\} \\ &= \Pr\{(\mathbf{X}_{iA}\boldsymbol{\beta} + u_{iA}) \geq E[U_i(SQ)]\} \end{aligned}$$

Digression II: The Random Utility Model

Normalize:

$$E[U_i(SQ)] = 0$$

Then:

$$\begin{aligned}\Pr(Y = A) &= \Pr\{(\mathbf{X}_{iA}\boldsymbol{\beta} + u_{iA}) \geq 0\} \\ &= \Pr\{u_{iA} \geq -\mathbf{X}_{iA}\boldsymbol{\beta}\} \\ &= F(\mathbf{X}_{iA}\boldsymbol{\beta})\end{aligned}$$

Another Model: Complementary Log-Log

Uses:

$$\Pr(Y_i = 1) = 1 - \exp[-\exp(\mathbf{X}_i\beta)]$$


or

$$\ln\{-\ln[1 - \Pr(Y_i = 1)]\} = \mathbf{X}_i\beta$$

Likelihood is:

$$\begin{aligned} \ln L = & \sum_{i=1}^N Y_i \ln\{1 - \exp[-\exp(\mathbf{X}_i\beta)]\} + \\ & (1 - Y_i) \ln\{1 - \{1 - \exp[-\exp(\mathbf{X}_i\beta)]\}\} \end{aligned}$$

Logit and C-log-log CDFs



All require that:

- “Threshold” = $Y^* > 0$
- $E(u_i | \mathbf{X}, \beta) = 0$
- $\text{Var}(u_i) = \frac{\pi^2}{3}$ or 1.0.

In general:


- The Universe: Logit $>$ _{Probit}
- The (Social Science) Universe: Meh...
- $\hat{\beta}_{\text{Logit}} \approx 1.8 \times \hat{\beta}_{\text{Probit}}$
- Four reasons to prefer / use logit

A Toy Example

```
> set.seed(7222009)
> ystar<-rnorm(100,0.5,0.5)
> y<-ifelse(ystar>0.5,1,0)
> x<-ystar+(0.5*rnorm(100))
> data<-data.frame(ystar,y,x)
```

```
> head(data)
      ystar y      x
1  0.17977 0  0.2677
2  0.79428 1  1.5079
3  0.82408 1  0.8842
4  0.24658 0  0.8172
5  0.50966 1  1.1255
6 -0.07852 0 -0.6506
```

A Toy Example



Model Comparisons

```
> mylogit<-glm(y~x,family=binomial(link="logit"), data=data)
> myprobit<-glm(y~x,family=binomial(link="probit"), data=data)
> mycloglog<-glm(y~x,family=binomial(link="cloglog"), data=data)
```

Logit, Probit, and C-Log-Log Models (Simulated Data)

	<i>Logit</i>	<i>Probit</i>	<i>C-Log-Log</i>
<i>X</i>	2.428*** (0.500)	1.458*** (0.272)	1.613*** (0.309)
Constant	-0.861*** (0.318)	-0.519*** (0.183)	-1.048*** (0.250)
Observations	100	100	100
Log Likelihood	-49.690	-49.490	-49.522
Akaike Inf. Crit.	103.380	102.979	103.044

Note:


* $p < 0.1$; ** $p < 0.05$; *** $p < 0.01$

Comparing Models (continued)

Note:

- z s, P s, $\ln L$ s, AICs nearly identical
- Residuals, too
- $\hat{\beta}_{\text{Logit}}$ is $\frac{2.428}{1.458} = 1.54 \times \hat{\beta}_{\text{Probit}}$

Toy Example: Predicted Probabilities



Note: C-Log-Log Isn't "Reversible"

Suppose we generate a new dependent variable:

$$Y_{i\text{New}} = 1 - Y_i$$

What happens to our estimates?

	$\hat{\beta}_0$			$\hat{\beta}_1$		
	Y		Y_{New}	Y		Y_{New}
Probit	-0.52	\leftrightarrow	0.52	1.46	\leftrightarrow	-1.46
Logit	-0.86	\leftrightarrow	0.86	2.43	\leftrightarrow	-2.43
C-Log-Log	-1.05	\leftrightarrow	0.11	1.61	\leftrightarrow	-1.66